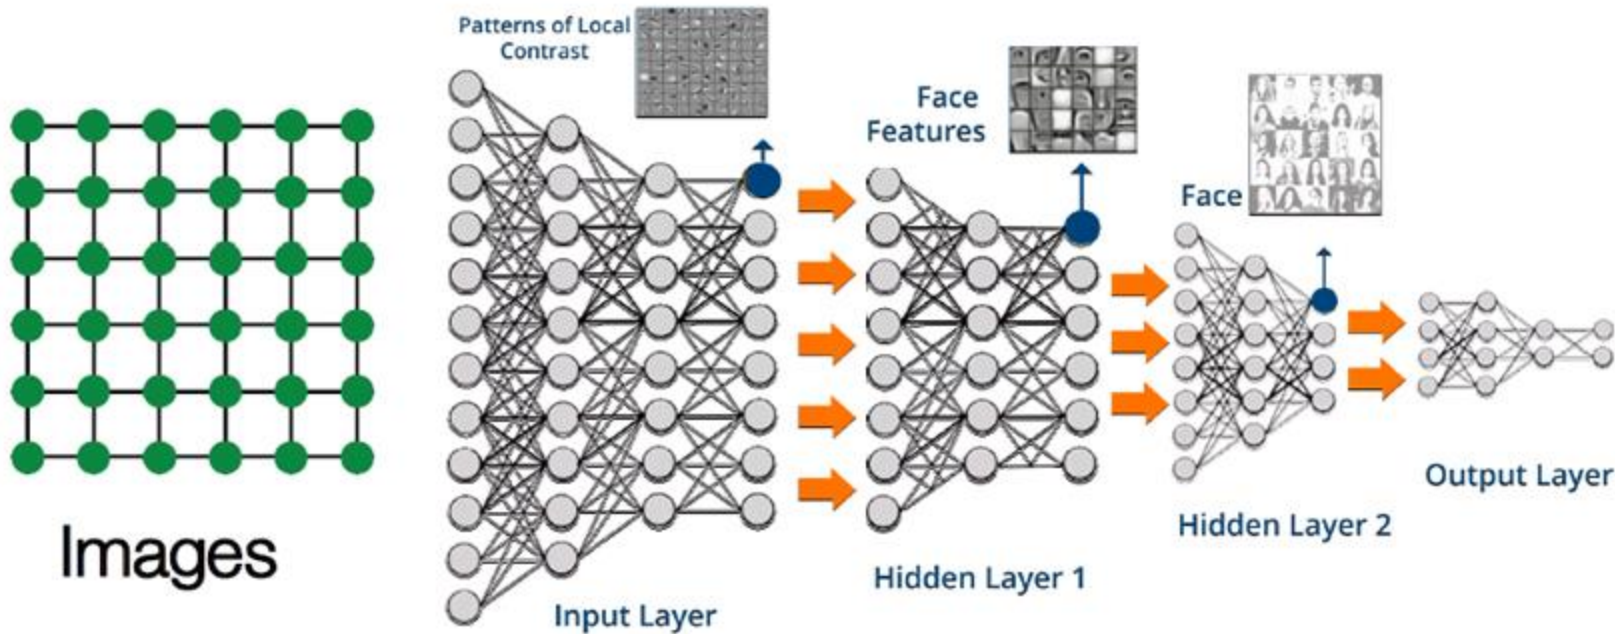



Graph Neural Network

Jure Leskovec, Michele Catasta, Stanford University

Modern DL Toolbox is Designed for Sequences & Grids




Text/Speech




From Images to Graphs

Single CNN layer with 3x3 filter:



Image




Graph


Transform information at the neighbors and combine it:

- Transform “messages” h_i from neighbors: $W_i h_i$
- Add them up: $\sum_i W_i h_i$

How about these graphs?



or this:



Biological networks, Medical networks, Social networks, Information networks,
Knowledge graphs, Communication networks, Web graph

Problem Setup

- **Encoder:** Map a node to a low-dimensional vector:

$$\text{ENC}(v) = \mathbf{z}_v$$

node in the input graph \rightarrow v


\mathbf{z}_v \leftarrow d-dimensional embedding

- **Similarity function** defines how relationships in the input network map to relationships in the embedding space:

$$\text{similarity}(u, v) \approx \mathbf{z}_v^\top \mathbf{z}_u$$

Similarity of u and v in the network \rightarrow $\text{similarity}(u, v)$

$\mathbf{z}_v^\top \mathbf{z}_u$ \leftarrow dot product between node embeddings



Shallow Encoding

- Simplest encoding approach: **encoder is just an embedding-lookup**

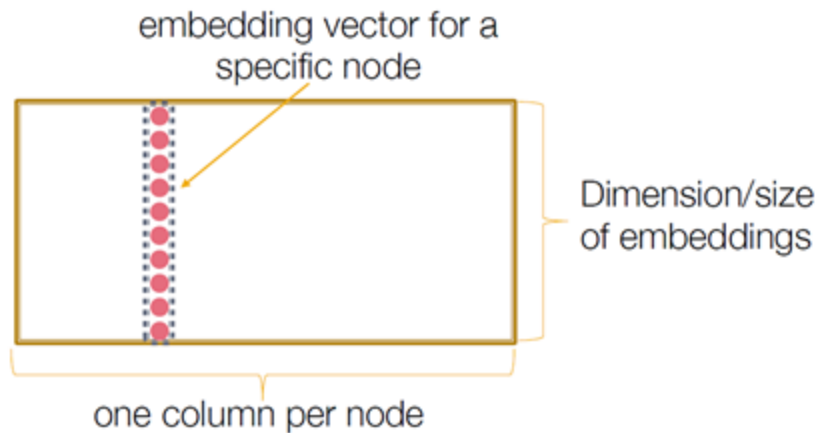
$$\text{ENC}(v) = \mathbf{Z}\mathbf{v}$$

$\mathbf{Z} \in \mathbb{R}^{d \times |\mathcal{V}|}$ matrix, each column is a node embedding [what we learn!]

$\mathbf{v} \in \mathbb{I}^{|\mathcal{V}|}$ indicator vector, all zeroes except a one in column indicating node v

embedding matrix

$\mathbf{Z} =$




Each node is assigned to a unique embedding vector

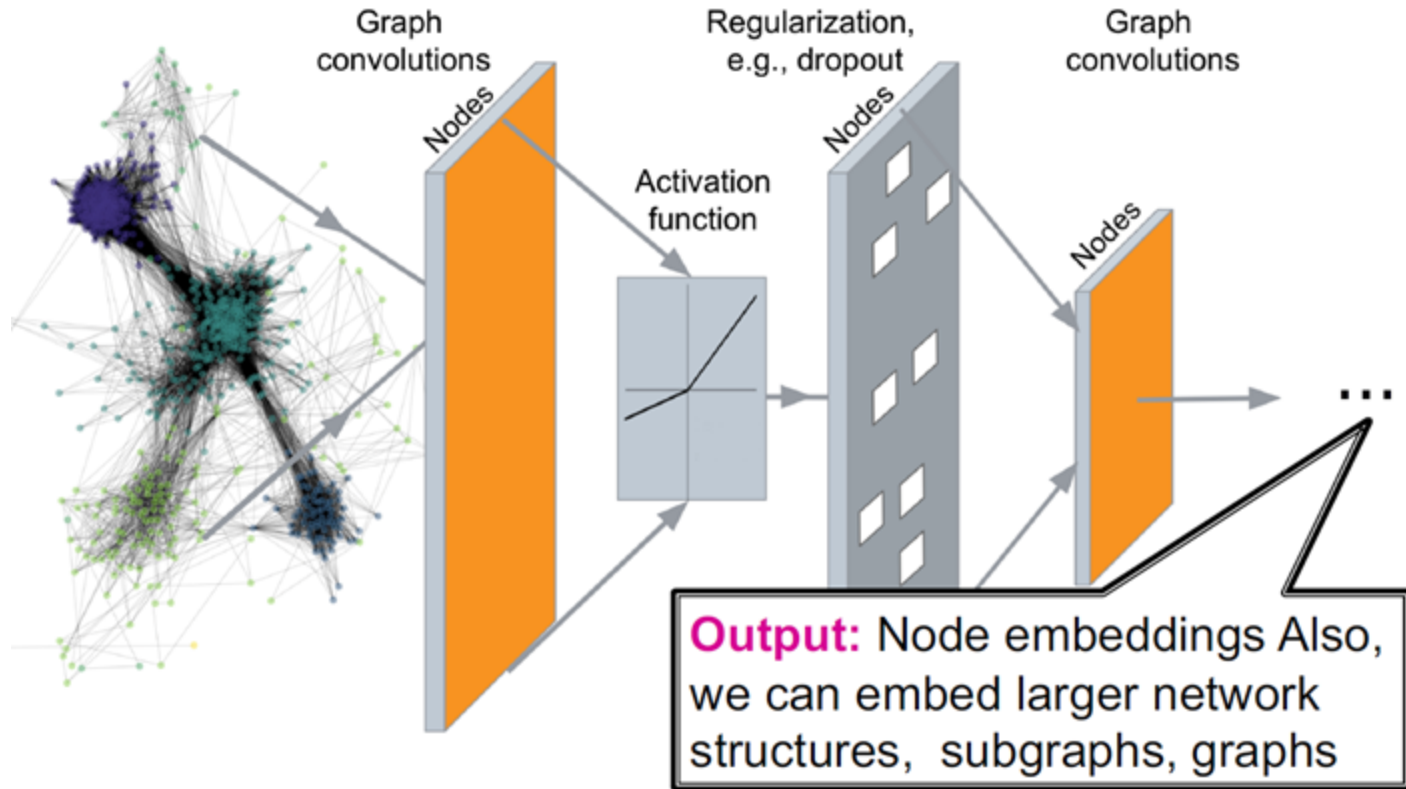
Many methods: DeepWalk, node2vec, TransE

Limitations of Shallow Encoding

- **$O(|V|)$ parameters are needed:**
 - No sharing of parameters between nodes
 - Every node has its own unique embedding
- **Inherently “transductive”:**
 - Cannot generate embeddings for nodes that are not seen during training
- **Do not incorporate node features:**
 - Many graphs have features that we can and should leverage

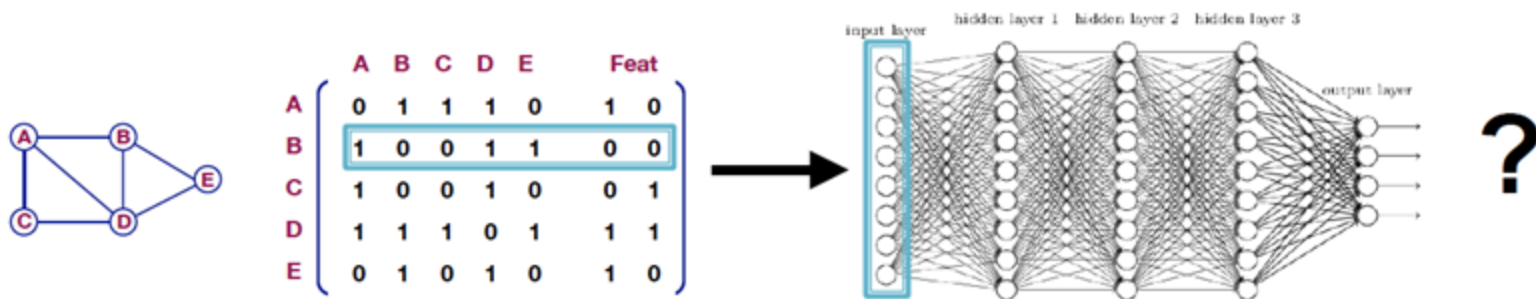


Deep Graph Encoders



A Naive Approach

- Join adjacency matrix and features
- Feed them into a deep neural net:




- Issues with this idea:
 - $O(N)$ parameters
 - Not applicable to graphs of different sizes
 - Not invariant to node ordering

Setup


- Assume we have a graph G :
 - V is the **vertex set**
 - A is the **adjacency matrix** (assume binary)
 - $X \in \mathbb{R}^{m \times |V|}$ is a matrix of **node features**
 - Node features:
 - Social networks: User profile, User image
 - Biological networks: Gene expression profiles, gene functional information
 - No features:
 - Indicator vectors (one-hot encoding of a node)
 - Vector of constant 1: $[1, 1, \dots, 1]$

Message Passing of Nodes

Idea: Node's neighborhood defines a computation graph



Determine node
computation graph



Propagate and
transform information

Advantages:



Capturing the Structure

Utilizing feature information

Learn how to propagate information across the
graph to compute node features


Aggregate Neighbors

- **Key idea:** Generate node embeddings based on **local network neighborhoods**



Shared Params of Neural Networks

- **Intuition:** Nodes aggregate information from their neighbors using neural networks




Aggregate Neighbors

- **Intuition:** Network neighborhood defines a computation graph

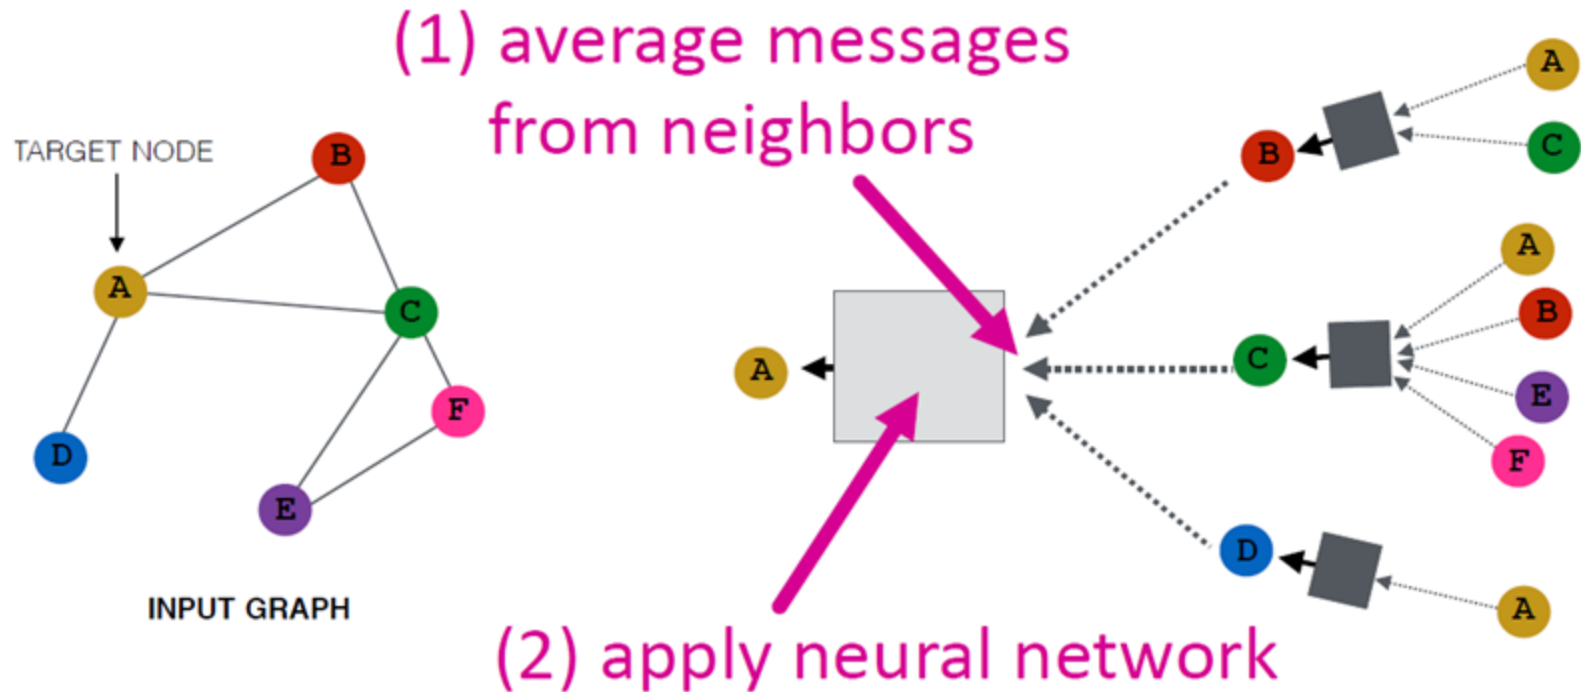
Every node defines a computation graph based on its neighborhood!



INPUT GRAPH



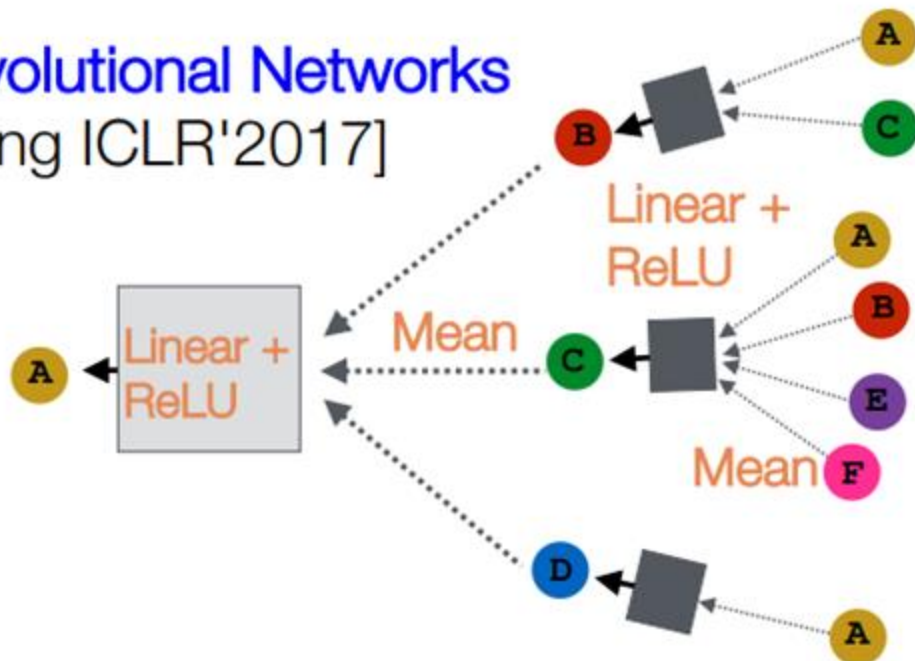

- **Basic approach:** Average information from neighbors and apply a neural network



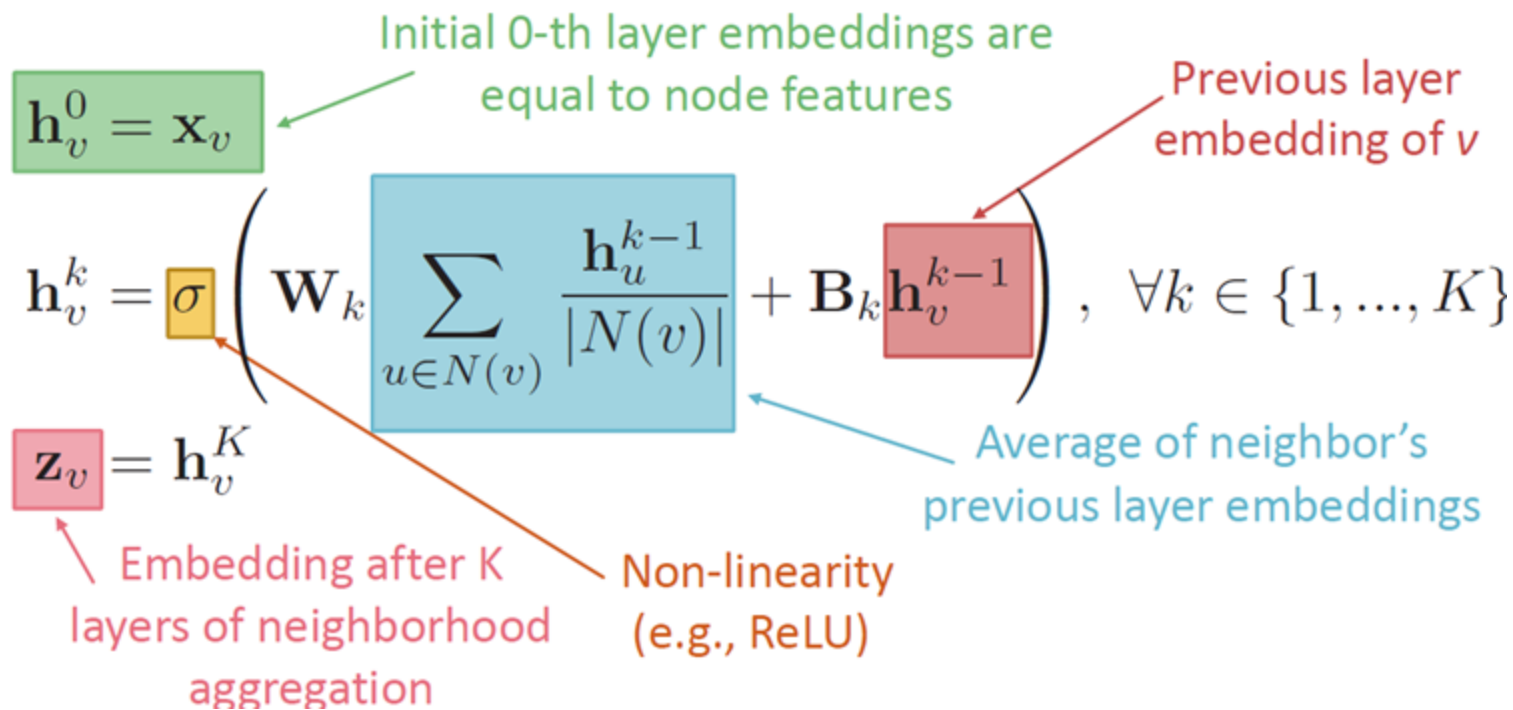
- Many model variants have been proposed with different choice of neural networks.

Scarselli et al., 2009b; Battaglia et al., 2016; Defferrard et al., 2016; Duvenaud et al., 2015; Hamilton et al., 2017a; Kearnes et al., 2016; Kipf & Welling, 2017; Lei et al., 2017; Li et al., 2016; Velickovic et al., 2018; Verma & Zhang, 2018; Ying et al., 2018; Zhang et al., 2018

Graph Convolutional Networks [Kipf & Welling ICLR'2017]



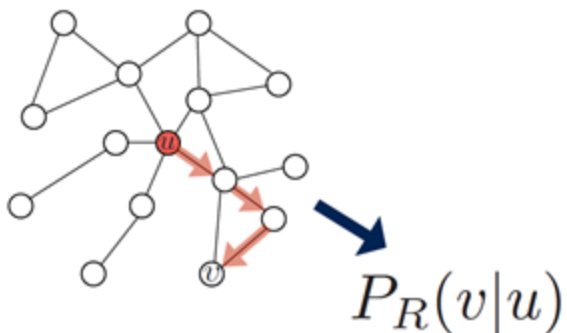
- **Basic approach:** Average neighbor messages and apply a neural network




Unsupervised Training

Similar nodes should have similar embeddings

Can define a loss function based on results from node2vec, deepwalk, struc2vec




Random Walk



Node2vec: biased random walk

Supervised Training

Node Classification



E.g., a drug-drug interaction network

Directly train the model for a supervised task
(e.g., **node classification**)

$$\mathcal{L} = \sum_{v \in V} y_v \log(\sigma(\mathbf{z}_v^\top \boldsymbol{\theta})) + (1 - y_v) \log(1 - \sigma(\mathbf{z}_v^\top \boldsymbol{\theta}))$$

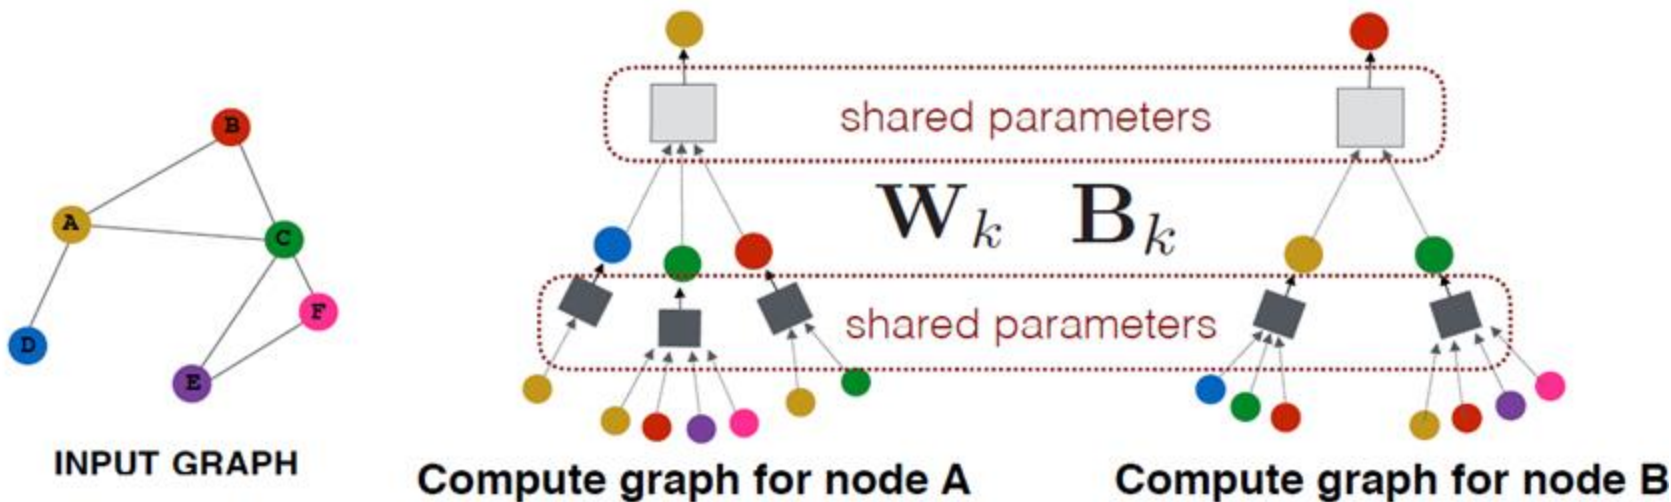
Encoder output:
node embedding

Classification
weights

Node class
label

Inductive Capability

- The same aggregation parameters are shared for all nodes:
 - The number of model parameters is sublinear in $|V|$ and we can **generalize to unseen nodes!**



Application: Pinterest

Human curated collection of pins



Very ape blue structured coat

Hilly Giffly

Picked for you
Street style



Hans Wegner chair

Robert and Board

Promoted by
Robert & Board



This is just a beautiful image for thoughts. Yay or nay, your choice.

Annie Tong
Plantation

Pin: A visual bookmark someone has saved from the internet to a board they've created.

Pin: Image, text, link

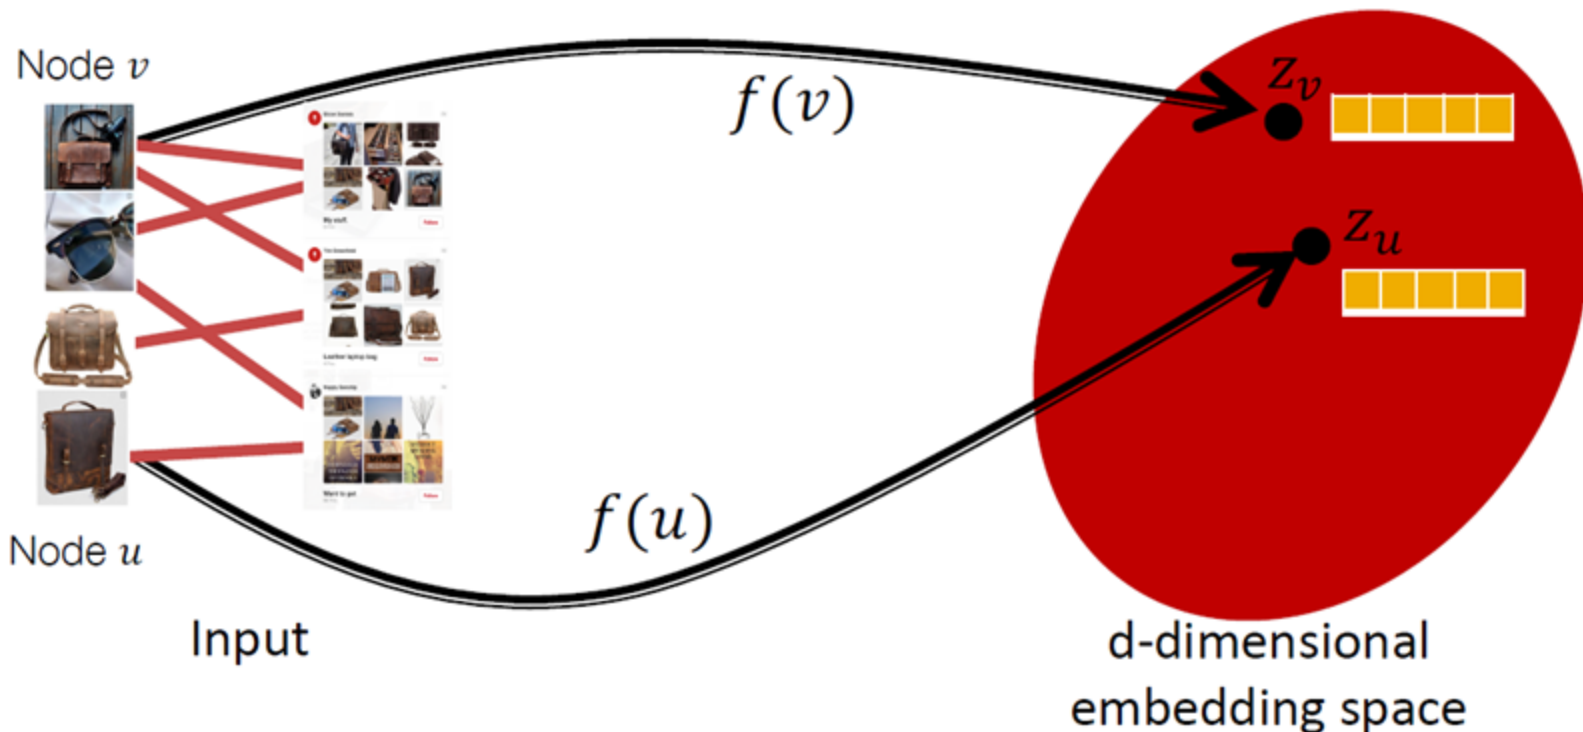


Board: A collection of ideas (pins having something in common)

- **PinSage** graph convolutional network:
 - **Goal:** Generate embeddings for nodes (e.g., Pins/images) in a web-scale Pinterest graph containing billions of objects
 - **Key Idea:** Borrow information from nearby nodes
 - E.g., bed rail Pin might look like a garden fence, but gates and beds are rarely adjacent in the graph




- Pin embeddings are essential to various tasks like recommendation of Pins, classification, clustering, ranking
 - Services like “Related Pins”, “Search”, “Shopping”, “Ads”



Goal: Map nodes to d-dimensional embeddings such that nodes that are related are embedded close together

Task: Recommend related pins to users



Task: Learn node embeddings z_i such that

$$d(z_{cake1}, z_{cake2}) < d(z_{cake1}, z_{sweater})$$

■ Challenges:

- **Massive size:** 3 billion nodes, 20 billion edges
- **Heterogeneous data:** Rich image and text features

Goal: Identify target pin among 3B pins

- **Issue:** Need to learn with resolution of 100 vs. 3B
- **Idea:** Use harder and harder negative samples
- Include more and more hard negative samples for each epoch



Source pin



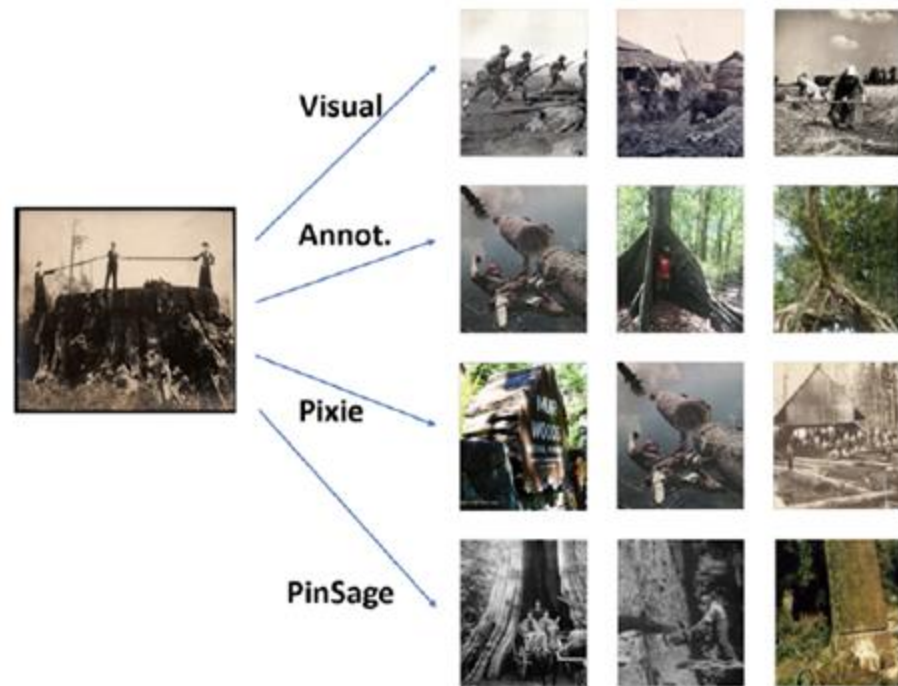

Positive



Easy negative



Hard negative



Pixie: Random Walk (Dksombatchai et al., 2018)

Further Reading

PyTorch
Geometric

Tutorials and overviews:

- Relational inductive biases and graph networks (Battaglia et al., 2018)
- Representation learning on graphs: Methods and applications (Hamilton et al., 2017)

Attention-based neighborhood aggregation:

- Graph attention networks (Hoshen, 2017; Velickovic et al., 2018; Liu et al., 2018)

Embedding entire graphs:

- Graph neural nets with edge embeddings (Battaglia et al., 2016; Gilmer et al., 2017)
- Embedding entire graphs (Duvenaud et al., 2015; Dai et al., 2016; Li et al., 2018) and graph pooling (Ying et al., 2018, Zhang et al., 2018)
- Graph generation and relational inference (You et al., 2018; Kipf et al., 2018)
- How powerful are graph neural networks (Xu et al., 2017)

Embedding nodes:

- Varying neighborhood: Jumping knowledge networks (Xu et al., 2018), GeniePath (Liu et al., 2018)
- Position-aware GNN (You et al. 2019)

Spectral approaches to graph neural networks:

- Spectral graph CNN & ChebNet (Bruna et al., 2015; Defferrard et al., 2016)
- Geometric deep learning (Bronstein et al., 2017; Monti et al., 2017)

Other GNN techniques:

- Pre-training Graph Neural Networks (Hu et al., 2019)
- GNNExplainer: Generating Explanations for Graph Neural Networks (Ying et al., 2019)